

Cyclical Dynamics in Idiosyncratic Labor Market Risk

Kjetil Storesletten

University of Oslo and Centre for Economic Policy Research

Chris I. Telmer

Carnegie Mellon University

Amir Yaron

University of Pennsylvania and National Bureau of Economic Research

Is individual labor income more risky in recessions? This is a difficult question to answer because existing panel data sets are so short. To address this problem, we develop a generalized method of moments estimator that conditions on the macroeconomic history that each member of the panel has experienced. Variation in the cross-sectional variance between households with differing macroeconomic histories allows us to incorporate business cycle information dating back to 1930, even though our data do not begin until 1968. We implement this estimator using household-level labor earnings data from the Panel Study of Income Dynamics. We estimate that idiosyncratic risk is (i) highly persistent, with an annual autocorrelation coefficient of 0.95, and (ii) strongly countercyclical, with a conditional standard

This paper circulated previously as a component of the paper "Asset Pricing with Idiosyncratic Risk and Overlapping Generations." In addition to participants at seminars and conferences, we thank Dave Backus, Geert Bekaert, David Bowman, Martin Browning, John Cochrane, Wouter den Haan, Mick Devereux, Ron Gallant, René Garcia, Rick Green, John Heaton, Burton Hollifield, Mark Huggett, Narayana Kocherlakota, Per Krusell, Deborah Lucas, Pierre Mella-Barral, Bob Miller, Michael Parkin, Victor Ríos, Amlan Roy, Tony Smith, Paul Willen, Steve Zeldes, Stan Zin, and three anonymous referees for helpful comments and suggestions. We have benefited from the support of National Science Foundation grant SES-9987602 and the Rodney White Center at Wharton. Valuable research assistance was provided by Xiaorong Dong.

[*Journal of Political Economy*, 2004, vol. 112, no. 3]
© 2004 by The University of Chicago. All rights reserved. 0022-3808/2004/11203-0007\$10.00

deviation that increases by 75 percent (from 0.12 to 0.21) as the macroeconomy moves from peak to trough.

I. Introduction

The interaction between cross-sectional risk and aggregate risk plays an increasingly important role in macroeconomics and finance. For example, Bernanke and Gertler (1989), Carlstrom and Fuerst (1997), Cooley, Marimon, and Quadrini (2002), and others argue that financing constraints are an important propagation mechanism because they generate asymmetric effects of aggregate shocks on the cross-sectional distribution of firms. Davis and Haltiwanger (1992), Caballero and Hamour (1994), Boeri (1996), Foote (1998), and others examine cyclical dynamics in job creation and destruction. Rampini (2000) argues that incentive constraints associated with entrepreneurial activity give rise to countercyclical cross-sectional risk. Krusell and Smith (1999), Storesletten, Telmer, and Yaron (2001*b*), and Krebs (2002, 2003) ask if countercyclical idiosyncratic risk is important for the welfare costs of business cycles. Finally, and most closely related to this paper, Lucas (1994), Heaton and Lucas (1996), Krusell and Smith (1997), Marcet and Singleton (1999), Balduzzi and Yao (2000), Alvarez and Jermann (2001), Lustig (2001), Storesletten et al. (2001*a*), Brav, Constantinides, and Geczy (2002), Cogley (2002), and Sarkissian (2003) follow Mankiw (1986) and Constantinides and Duffie (1996) and study models in which assets with higher expected returns do badly at times of higher cross-sectional labor income risk.

The size of cyclical variation in the cross-sectional distribution of labor income risk is thus an important question. It is hard to answer, however, because panel data sets are so short. Even the Panel Study of Income Dynamics (PSID)—the panel with the longest time dimension—covers only four or five business cycles. We address this problem. We incorporate macroeconomic information dating back to 1930, in spite of the fact that our panel data do not begin until 1968. To understand what we do, consider two “cohorts” of individuals, the first born in 1910 and the second in 1930. Everyone is subject to idiosyncratic labor market shocks, some fraction of which are highly persistent. The conditional variance of these shocks is countercyclical, increasing during contractions and decreasing during expansions. Suppose that we have labor income data on each cohort when its members are 60 years old. That is, we have 1970 data on the 1910 cohort and 1990 data on the 1930 cohort. What we shall observe—given the high persistence and the countercyclical idiosyncratic risk—is more cross-sectional dispersion among the 1910 cohort than among the 1930 cohort. The reason is that the

former worked through more contractionary years than the latter, including the Great Depression. More generally, we shall see variation in the cross-sectional variance between any two cohorts of similar ages who have worked through different macroeconomic histories. This is the essence of our procedure. Even though the time dimension of our panel data is limited to 1968–93, we have a rich cross section of *ages* in each year of the panel. We can therefore use macroeconomic data to characterize the working history of each household in the panel and then use cross-sectional variation between cohorts of similar ages to identify the cyclical idiosyncratic-risk effects we are after.

More specifically, we focus on the properties of household-level labor market earnings from the PSID. We model idiosyncratic risk as a class of ARMA(1,1) processes with a regime-switching component in the conditional variance. The regime is identified by using macroeconomic data to classify each year between 1930 and 1993 as either a contraction or an expansion. A critical feature is finiteness: we make strong assumptions about initial conditions that allow us to base a generalized method of moments (GMM) estimator on age-dependent moments.

We find robust evidence that idiosyncratic earnings risk is both highly persistent and countercyclical. Estimates of annual autocorrelation range from 0.94 to 0.96. Estimates of conditional standard deviations increase by roughly 75 percent (from 0.12 to 0.21) as the macroeconomy moves from expansion to contraction. We use graphical methods to make these GMM estimates transparent. We show that high autocorrelation is driven by a linearly increasing pattern of cross-sectional variance with age.¹ Countercyclical volatility is driven by cohort effects in the cross-sectional variance, something that is clearly evident in a simple scatter plot.

The remainder of the paper is organized as follows. Section II presents our time-series model, and Section III discusses our data. Section IV presents a graphical analysis that shows that most of our results can be understood visually. Section V presents GMM estimates of the model in Section II. Section VI reconciles our estimates of variation in conditional variance with variation in the overall (age-pooled) cross section. Section VII presents conclusions.

II. Time-Series Model for Idiosyncratic Risk

We begin with a time-series model for idiosyncratic labor earnings risk. Its most distinctive property—which will be emphasized in estimation—is that it depends explicitly on age. We denote Y_t as the logarithm of

¹ Deaton and Paxson (1994) have also documented this evidence using a different data source, the Consumer Expenditure Survey.

per capita labor earnings at time t and y_{it}^h as the logarithm of labor earnings for household i of age h at time t . Log earnings y_{it}^h is decomposed into two components:

$$y_{it}^h = g(x_{it}^h, Y_t) + u_{it}^h \quad (1)$$

The component $g(x_{it}^h, Y_t)$ incorporates aggregate earnings as well as x_{it}^h deterministic components of household-specific earnings attributable to age, education, and so on.² The component u_{it}^h is the random component of a household's earnings that is idiosyncratic to them. It is identified by $E(u_{it}^h | Y_t, x_{it}^h) \equiv \tilde{E}_t(u_{it}^h) = 0$, for all dates t , where \tilde{E}_t denotes the cross-sectional mean.

The specification of $g(x_{it}^h, Y_t)$ is critical for identifying u_{it}^h , but beyond that, it is not central. We defer its discussion until Section V. For u_{it}^h , we use an ARMA(1,1) with a regime-switching conditional variance:

$$\begin{aligned} u_{it}^h &= \alpha_i + z_{it}^h + \epsilon_{it} \\ z_{it}^h &= \rho z_{i,t-1}^h + \eta_{it} \end{aligned} \quad (2)$$

where $\alpha_i \sim \text{Niid}(0, \sigma_\alpha^2)$, $\epsilon_{it} \sim \text{Niid}(0, \sigma_\epsilon^2)$, $\eta_{it} \sim \text{Niid}(0, \sigma_\eta^2)$,

$$\sigma_i^2 = \begin{cases} \sigma_E^2 & \text{if aggregate expansion at date } t \\ \sigma_C^2 & \text{if aggregate contraction at date } t, \end{cases}$$

and $z_{it}^0 = 0$. The variable α_i is a "fixed effect": a shock received once at birth and then retained throughout life (hence the absence of a time subscript). The variables z_{it}^h and ϵ_{it} are persistent and transitory shocks, respectively (we make age h explicit only when the conditional distribution of a variable depends on it). What we mean by "countercyclical volatility" is that $\sigma_C > \sigma_E$. We do not impose this a priori.

We include fixed effects and transitory shocks so that we can better measure the primary objects of interest, σ_i^2 and ρ . For example, the way we shall identify these parameters involves how the cross-sectional variance of u_{it}^h increases with age. However, much of the overall variation in u_{it}^h is common to households of all ages. Were we to exclude fixed effects α_i , we would overstate the magnitude of σ_i^2 and, therefore, overstate the amount of idiosyncratic risk that is relevant for economic decision making. Along similar lines, the transitory shocks ϵ_{it} are included to mitigate the extent to which measurement error (and "true" transitory shocks) overstates our estimate of σ_i^2 .

The essence of our approach lies in the initial conditions: $z_{it}^0 = 0$ and $\alpha_i \sim \text{Niid}(0, \sigma_\alpha^2)$. These are strong assumptions. They rule out, for example, time variation in the distribution of α_i and dependence between

² In treating education attainment deterministically, we are probably underestimating the overall risks agents face. However, adding schooling decisions is beyond the scope of this paper.

α_i and, for a given i , subsequent realizations of z_{it}^h and ϵ_{it} . The benefits, however, are substantial. The initial conditions allow us to interpret equations (2) as a collection of finite processes and therefore allow us to condition on age. The implications are most apparent in the cross-sectional variances, for each age h , that will underly our GMM estimator:

$$\tilde{\text{Var}}(u_{it}^h) = \sigma_\alpha^2 + \sigma_\epsilon^2 + \sum_{j=0}^{h-1} \rho^{2j} [I_{t-j} \sigma_E^2 + (1 - I_{t-j}) \sigma_C^2], \quad (3)$$

where $\tilde{\text{Var}}$ denotes the cross-sectional variance, $I_t = 1$ if the economy is in an expansion at date t , and $I_t = 0$ otherwise. If $|\rho| < 1$, the summation term converges to the familiar $\sigma^2/(1 - \rho^2)$, where σ^2 is a probability weighted average of σ_E^2 and σ_C^2 . For finite h , however, $\tilde{\text{Var}}(u_{it}^h)$ is increasing in h , at a rate determined by the magnitude of ρ . This, in conjunction with a (roughly) linearly increasing age profile of sample moments for $\tilde{\text{Var}}(u_{it}^h)$, is what drives the relatively large estimate of ρ we shall arrive at in Section V.

Age dependence in equation (3) is also critical for how we assess countercyclical volatility. Consider, for instance, the cohort of 60-year-old workers in the first year of our panel, 1968. According to equation (3), the cross-sectional variance for this cohort involves indicator variables I_{t-j} dating as far back as 1930 (we assume that $h = 1$ corresponds to a 23-year-old worker). By classifying each year between 1930 and 1968 as either an expansion or a contraction, we can form history-dependent cross-sectional moments based on (3) that incorporate aggregate shocks dating back to 1930, far beyond the 1968–93 confines of our panel. Doing so across all 26 years of our panel is what identifies σ_E and σ_C . To see this, compare the 60-year-old workers in 1968 to the 60-year-old workers in 1993. Each has worked for 38 years. However, those in the 1968 cohort have worked during more contractions, including the Great Depression. The process (2) predicts that, if ρ is sufficiently large and $\sigma_C > \sigma_E$, then the cross-sectional variance of a particular cohort will be increasing in the number of contractionary years that its members worked through.³ Therefore, inequality among the 1968 cohort should exceed that of the 1993 cohort. This is exactly what we document below, and (anecdotally speaking) it is what drives our results on countercyclical volatility.

In summary, we are primarily interested in the autocorrelation, ρ , and countercyclical volatility, $\sigma_C > \sigma_E$. What will be pivotal for the former is how sample analogues of $\tilde{\text{Var}}(u_{it}^h)$ increase in h . This will also determine

³ It is not enough for ρ to simply be positive. Consider, e.g., two cohorts of the same age in which one worked through many contractionary years early in life and the other worked through only one contraction, but late in life. If ρ is sufficiently small, then the cross-sectional variance among the latter can be greater than among the former.

the *average* of σ_C and σ_E . What will be pivotal for the latter is variation in the cross-sectional variation between households of similar ages that have worked through different macroeconomic histories.

III. PSID Data

We use annual PSID data, 1968–93, defined at the household level. We define “earnings” as wage receipts of all adult household members, plus any transfers received such as unemployment insurance, worker’s compensation, transfers from nonhousehold family members, and so on. Transfers are included because we wish to measure idiosyncratic risk *net* of the implicit insurance mechanisms that these payments often represent. Along similar lines, we study the household as a single unit so as to abstract from shocks that are insured via intrahousehold variation in labor force participation.

We depart from the common practice of using a longitudinal panel: an equal number of time-series observations on a fixed cross section of households. A longitudinal panel necessarily features an average age that increases with time. This is problematic for our approach, which emphasizes restrictions between age and aggregate variation. For instance, were we to use a longitudinal panel for the years 1968–93, a large fraction of the household heads will have been retired, or at least been in their late earning years, during the last two of only five business cycles witnessed during 1968–93.⁴ A longitudinal panel will also feature a relatively small sample size and be subject to survivorship bias in that only the relatively stable households are likely to remain in the panel for all 26 years.

We overcome these issues by constructing a sequence of overlapping three-year subpanels. For each of the years 1968–93, we construct a three-year subpanel consisting of households that reported strictly positive total household earnings (inclusive of transfers) for the given year and the next two consecutive years in the survey. For example, our 1971 subpanel is essentially a longitudinal panel of 2,036 households over the years 1971, 1972, and 1973. Doing this for all years results in a sequence of 24 overlapping subpanels. The overall data structure contains more than enough time-series information to identify the parameters in equation (2), while at the same time survivorship bias is mitigated and the cross-sectional distribution of age is quite stable. The mean and standard deviation of the average age in each subpanel are 40.8 and 1.0, respectively. The mean and the standard deviation (across panels) of the number of households are 2,449 and 322, respectively.

⁴ Specifically, average age would increase from 39 to 64 over the years 1968–93.

Additional details, including a number of filters related to measurement error and demographic stability, are outlined in the Appendix.

IV. Graphical Analysis

Most of our econometric results can be anticipated by simply looking at the data. The important dimensions are variations in the cross-sectional variance according to age and according to time. Loosely speaking, variation across age is what drives our estimate of autocorrelation, ρ , whereas variation across time, and how it relates to age, is what drives our estimates of σ_E and σ_C .

To show this, we use a more reduced-form representation of the data—a dummy variable regression similar to that in Deaton and Paxson (1994)—and ask what it suggests about our model's parameters. Denoting the cross-sectional variances in equation (3) as $\sigma_{h,t}^2 \equiv \tilde{\text{Var}}(u_{it}^h)$, we decompose their variance into cohort and age effects:

$$\sigma_{h,t}^2 = a_c + b_h + e_{h,t} \quad (4)$$

where $c = t - h$ denotes a cohort (i.e., birth year).⁵ The parameters a_c and b_h are cohort and age effects, and $e_{h,t}$ are residuals. Our model represents restrictions on the age/time/cohort effects captured in the a_c and b_h parameters. We impose and test them in Section V. Here, we show what the unrestricted estimates imply about the parameters in (3).

We begin with age. Figure 1a plots the age effects, b_h , from the regression (4).⁶ The graph shows that the variance among the young is substantial and that it increases by a factor of 2.3 between ages 23 and 60. Moreover, the increase is approximately linear. Deaton and Paxson (1994) report similar results using data from the Consumer Expenditure Survey. Inspection of equation (3) suggests that the initial dispersion can identify $\sigma_\alpha^2 + \sigma_\epsilon^2$, the rate of increase can (loosely speaking) identify the average of σ_E and σ_C , and the linear shape can identify ρ . This suggests a value of ρ very close to unity. Section V formalizes this and shows that the near-unit root implications are robust to the incorporation of information not represented in the graph, most notably autocovariances.

Figure 1b plots the cohort coefficients, but in a different fashion. An implication of our model is that, when we control for age (as eq. [4] does), a cohort that has worked through more contractions than another

⁵ As is well known, an alternative decomposition is age and time effects. We choose cohort effects because, as fig. 1b will show, they are closely associated with the essence of our procedure. Moreover, Sec. VI deals with time effects, albeit in a somewhat different fashion.

⁶ We plot the variance instead of the standard deviation because it will clearly delineate the case of $\rho = 1$, which implies linearity for the former but not for the latter. The coefficients b_h are scaled so that the graph passes through the unconditional (across cohorts) variance of agents of age 40.

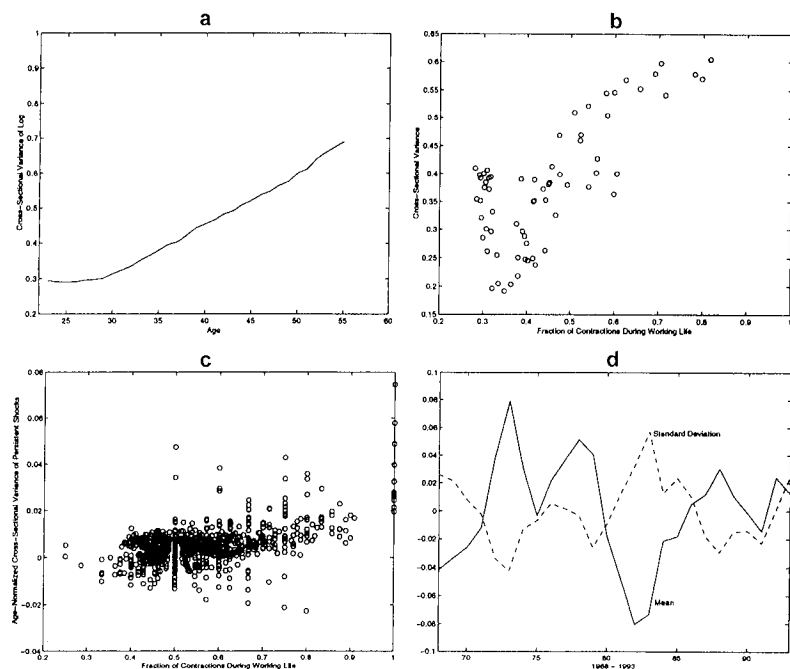


FIG. 1.—Cross-sectional variance of earnings, based on the idiosyncratic component (u_{it}^h from eq. [2]) of log labor earnings plus transfers from the PSID, 1968–93. *a*, Cross-sectional variance by age. The variances control for “cohort effects” by regressing cohort age-specific cross-sectional variances on cohort age dummy variables as in Deaton and Paxson (1994) and Storesletten et al. (in press). The points in the graph are the age coefficients, rescaled to match the level of variance at age 40. *b*, Macroeconomic history vs. cross-sectional variance (cohorts). The panel plots the cohort coefficients from the dummy variable regression underlying panel *a* against the fraction of contractionary years during which the oldest members of that cohort were of working age. *c*, Macroeconomic history vs. cross-sectional variance. The panel plots the age-normalized cross-sectional variance of the persistent shock underlying eq. (5) against the fraction of contractionary years during which members of that age panel-year worked through. *d*, Cross-sectional moments by time. The panel reports the linearly detrended cross-sectional mean of log income (y_{it}^h from eq. [1]) and the linearly detrended standard deviation of u_{it}^h , pooled across all ages, for each year 1968–93. The standard deviation is additively scaled for graphical reasons. The correlation coefficient between the two series is $-.74$. Panel *d* is robust to (i) alternative methods of detrending the mean and (ii) using the coefficient of variation instead of the standard deviation. Further details are given in Sec. III.

should, on average, exhibit greater cross-sectional dispersion (as long as ρ is large and $\sigma_{\epsilon} > \sigma_{\eta}$). Figure 1*b* shows that this is a characteristic of our panel. The cohort coefficients a_c are plotted against the fraction of years during the working life of each cohort that were contractions (defined as the National Income and Product Accounts measure of gross domestic product growth being below average for the elders of each cohort). Although there are a number of cohorts that worked through

relatively few contractions and nevertheless exhibit substantial cross-sectional dispersion (i.e., the cluster of points in the southwest corner of the graph), the positive relationship is apparent. The ordinary least squares slope coefficient is 0.94 with a standard error of 0.10. This relationship between cross-sectional variance and macroeconomic history is the main reason that our conclusions regarding variation in σ_t^2 differ from those of previous studies, in particular Heaton and Lucas (1996), who found relatively small effects.

In Figure 1c we move beyond the variance decomposition (4) and impose some structure from our model. Suppose that $\rho = 1$. Then the cross-sectional variances (3) imply

$$\begin{aligned}\sigma_{h,t}^2 &= \sigma_\alpha^2 + \sigma_\epsilon^2 + (h - n_{h,t})\sigma_E^2 + n_{h,t}\sigma_C^2 \\ \Rightarrow \frac{\sigma_{h,t}^2}{h} &= \frac{\sigma_\alpha^2 + \sigma_\epsilon^2}{h} + \sigma_E^2 + f_{h,t}(\sigma_C^2 - \sigma_E^2),\end{aligned}\quad (5)$$

where $n_{h,t}$ is the number of contractions that households of age h in panel year t have worked through, and $f_{h,t}$ expresses $n_{h,t}$ as a fraction of age h . We set $\sigma_\alpha^2 + \sigma_\epsilon^2 = 0.3$, using the intercept from figure 1a. Figure 1c is a scatter plot of sample moments of $h^{-1}(\sigma_{h,t}^2 - 0.3)$ versus $f_{h,t}$. Although the data are obviously much noisier than those in figure 1b—we are now essentially plotting the raw data from the regression (4)—a positive relationship is apparent. The ordinary least squares slope coefficient is 0.023, which, according to equation (5), is an estimate of countercyclical volatility ($\sigma_C^2 - \sigma_E^2$). This is close to the value that we shall estimate in Section V.

Figure 1d provides one last piece of evidence on countercyclical volatility. It pools the age effects highlighted in figure 1a and plots the year-by-year “pooled” cross-sectional standard deviation in the PSID. It also plots the cross-sectional mean. Both the mean and the standard deviation are linearly detrended. Even at this informal level we see striking evidence of countercyclical volatility. The correlation between the mean and the standard deviation is $-.74$. The magnitude of the changes, however, must be interpreted with caution. In Section VI we examine the pooled cross-sectional variance in more detail. We show that, because it is a close cousin of the unconditional variance, it will always understate (time) variation in what we are ultimately interested in: the conditional variance σ_t^2 .

V. Estimation

Recall that equation (1) decomposed log earnings into two components, $y_{it}^h = g(x_{it}^h, Y_t) + u_{it}^h$ where $g(x_{it}^h, Y_t)$ captures deterministic cross-sectional

TABLE 1
SUMMARY STATISTICS OF PANELS

	Constant	Age	Age ²	Age ³	Education	Family Size
Estimate	7.512	-.013	.147	-.019	.161	.070
Standard error	.087	.006	.015	.001	.011	.001

NOTE.—Entries are estimates of the parameters of the first-stage regression (eq. [6]). The regression also includes year dummy variables, which are not reported here. The regression R^2 is .23.

variation x_{it}^h and aggregate variation Y_t and u_{it}^h captures idiosyncratic cross-sectional variation. We specify the first component as

$$g(x_{it}^h, Y_t) = \theta_0 + \theta_1' \mathbf{D}(Y_t) + \theta_2' \mathbf{x}_{it}^h \quad (6)$$

where $\mathbf{D}(Y_t)$ is a vector of year dummy variables, $t = 1968, \dots, 1993$, and \mathbf{x}_{it}^h is a vector composed of age, age squared, age cubed, and educational attainment for household i of age h at date t . Educational attainment is measured as the number of school years completed by the household head. Table 1 reports estimates of θ . The estimates and fit are quite similar to those in existing studies, implying a concave earnings function in age and education (e.g., Hubbard, Skinner, and Zeldes 1994).

Equation (6) allows us to identify the idiosyncratic component of income, u_{it}^h , which follows the process (2). We estimate its parameters using GMM. For the sake of transparency, we begin with an exactly identified system that corresponds closely with the graphical analysis of the previous section. Afterward, we add overidentifying restrictions. The cross-sectional variances, equations (3), represent one variance for each age/time pair, (h, t) :

$$\tilde{E}_t \left[(u_{it}^h)^2 - (\sigma_\alpha^2 + \sigma_\epsilon^2) - \sum_{j=0}^{h-1} \rho^{2j} [I_{t-j} \sigma_E^2 + (1 - I_{t-j}) \sigma_C^2] \right] = 0 \quad (7)$$

for all h, t , where h runs from 23 to 60 and t runs from 1968 to 1993. We do not have sufficient data to formulate sample moments for each h, t pair. We therefore form age cells of width three years and interpret the age of a household in a particular cell as the midpoint. Furthermore, in order to mimic the graph in figure 1a (i.e., one variance per age), we aggregate the moments (7) over time and form sample analogues:

$$\frac{1}{T} \sum_{t=1}^T \frac{1}{N_{ht}} \sum_{i=1}^{N_{ht}} \left\{ (u_{it}^h)^2 - (\sigma_\alpha^2 + \sigma_\epsilon^2) - \sum_{j=0}^{h-1} \rho^{2j} [I_{t-j} \sigma_E^2 + (1 - I_{t-j}) \sigma_C^2] \right\} = 0, \quad (8)$$

where N_{ht} is the sample size of households aged h at time t , and h is now to be interpreted as an age cell.

Row A of table 2 reports exactly identified estimates of ρ , σ_C , σ_E , and

TABLE 2
IDIOSYNCRATIC EARNINGS PROCESS: PARAMETER ESTIMATES

	ρ	σ_C	σ_E	σ_ϵ	σ_α	p -Value
A. Exactly identified (no autocovariance)	.963 (.053)	.162 (.062)	.088 (.048)	.562 (.091)
B. Exactly identified (with autocovariances)	.957 (.058)	.163 (.034)	.094 (.042)	.351 (.063)	.448 (.094)	...
C. Overidentified, $H \times T$.952 (.020)	.211 (.034)	.125 (.044)	.255 (.021)	.378 (.057)	.899
D. Overidentified, NBER indicators	.943	.201 (.041)	.119 (.022)	.255 (.025)	.386 (.064)	.939
E. Overidentified, unemployment indicators	.938 (.011)	.246 (.026)	.138 (.034)	.257 (.020)	.308 (.062)	.459
F. Overidentified, CRSP value-weighted return indicators	.939 (.037)	.159 (.031)	.084 (.019)	.235 (.011)	.423 (.019)	.967

NOTE.—Entries are GMM estimates of the process (2), based on PSID earnings data 1968–93. The parameters σ_ϵ^2 and σ_α^2 denote the conditional standard deviation of the persistent component, conditional on an aggregate expansion, E , or contraction, C . Estimates in rows A, B, and C define expansion/contraction in terms of the growth rate of U.S. GNP being above/below its mean, those in row D are based on the NBER definitions, those in row E are based on U.S. unemployment data, and those in row F are based on equity index returns. Row A uses moments (7) with ages 25, 35, 45, and 55. Row B adds one more parameter and one more moment, the first autocovariance, eq. (10) for age 40. Rows C–F use an overidentified system by using eqq. (7)–(9) in addition to second-order autocovariances. Standard errors are computed using the White (1980) estimator and incorporate sampling uncertainty from the first-stage regression, eq. (1) (reported in table 1). The p -value pertains to the overidentifying test of the moment conditions. Further details, including all data sources, are available in the Appendix.

$\sqrt{\sigma_\alpha^2 + \sigma_\epsilon^2}$ based on equations (8), using ages 25, 35, 45, and 55. Since cross-sectional variances alone cannot distinguish σ_ϵ and σ_α , we report their sum. The indicator variables I_t are identified by classifying each year, 1930–93, as an expansion or contraction according to whether the growth rate in U.S. gross national product per capita was above or below its sample mean (alternative definitions are reported later). The Appendix provides specific details of all GMM estimators, including how we construct the weighting matrix to incorporate differing sample sizes N_{ht} , the overlapping structure of our panel, the first-stage regression, and several other issues.

The results in table 2 accord well with the graphs in figure 1. Much of the overall level of dispersion gets attributed to some combination of fixed effects, σ_α , and transitory shocks/measurement error, σ_ϵ . We estimate that $\sigma_\alpha^2 + \sigma_\epsilon^2 = 0.56^2$, which corresponds closely to the intercept of $0.55^2 = 0.30$ from figure 1*a*. For the persistent shock z_{it}^h , we estimate an autocorrelation of $\rho = 0.963$, slightly lower than what we inferred graphically. For the conditional volatilities, we estimate $\sigma_C = 0.162$ (contraction) and $\sigma_E = 0.088$ (expansion), which, on average, are consistent with the size of the increase in variance in figure 1*a*.⁷ Viewed on their

⁷ Specifically, the conditional variance in fig. 1*a* increases by 0.4 between ages 25 and 55. With a homoskedastic unit root process for the persistent shock, this implies a conditional variance of $0.4/(55 - 25) = 0.013$, or a standard deviation of 0.114. The average of our point estimates, σ_C^2 and σ_E^2 , is $(0.162^2 + 0.088^2)/2 = 0.017 = 0.13^2$.

own, the estimates indicate a substantial amount of countercyclical volatility: an increase of 84 percent from peak to trough.

Next, we identify σ_α from σ_ϵ using cross-sectional autocovariances analogous to the variances (7):

$$\tilde{E}_t \left[u_{it}^h u_{it-1}^{h-1} - \sigma_\alpha^2 - \rho \sum_{j=1}^{h-1} \rho^{2(j-1)} [I_{t-j} \sigma_E^2 + (1 - I_{t-j}) \sigma_C^2] \right] = 0 \quad (9)$$

for all h, t . Again, we aggregate these moments over t for each age cell h and then form sample counterparts analogous to equations (8). Row B of table 2 reports exactly identified estimates obtained by adding just one moment from equations (9) to those underlying row A. We use $h = 40$, but the results are not sensitive to the particular age. The estimates of primary interest, ρ , σ_C , and σ_E , change only marginally relative to row A. For the other parameters, the majority of the previously combined variance gets attributed to fixed effects: $\sigma_\alpha = 0.44$ and $\sigma_\epsilon = 0.35$ (these add up to roughly the same variance as in row A).

The estimates in rows A and B of table 2 are transparent in the sense that, given the graphical evidence of Section IV, it is clear which features of the data are driving them. The remainder of the table examines how robust these findings are once we add overidentifying restrictions and climb inside the usual black box of GMM. We use sample analogues of the (h, t) pair variances and first-order autocovariances in equations (7) and (9), along with a set of second-order autocovariances analogous to (9). That is, we eliminate the time aggregation in equations (8) (and rows A and B) and just use the primitive moments.⁸ Row C of table 2 uses the same definitions of "expansion" and "contraction" as rows A and B (based on GNP growth), whereas rows D, E, and F use definitions based on NBER business cycle dates, the unemployment rate, and the stock market. We see slightly lower autocorrelation and, with the exception of row F, higher conditional volatility for the persistent shock. The higher conditional volatility comes at the expense of the transitory and fixed-effect shocks, where the combined volatility drops from roughly 0.32 to the neighborhood of 0.20. Apparently, the additional autocovariances have reclassified an important part of the variance from transitory/measurement error shocks into persistent shocks. Finally, the extent to which the persistent shock conditional variance is countercyclical is relatively stable, ranging from a 70 percent increase (expansion to contraction) in rows C and D to a 90 percent increase in row F.

⁸ Specifically, rows C–F use variances and first- and second-order autocovariances for ages 25, 35, 45, and 55, over the years 1968–93. The ages were chosen to ensure a sample size of at least 100 households per (h, t) pair. This makes for 312 moments and 307 degrees of freedom.

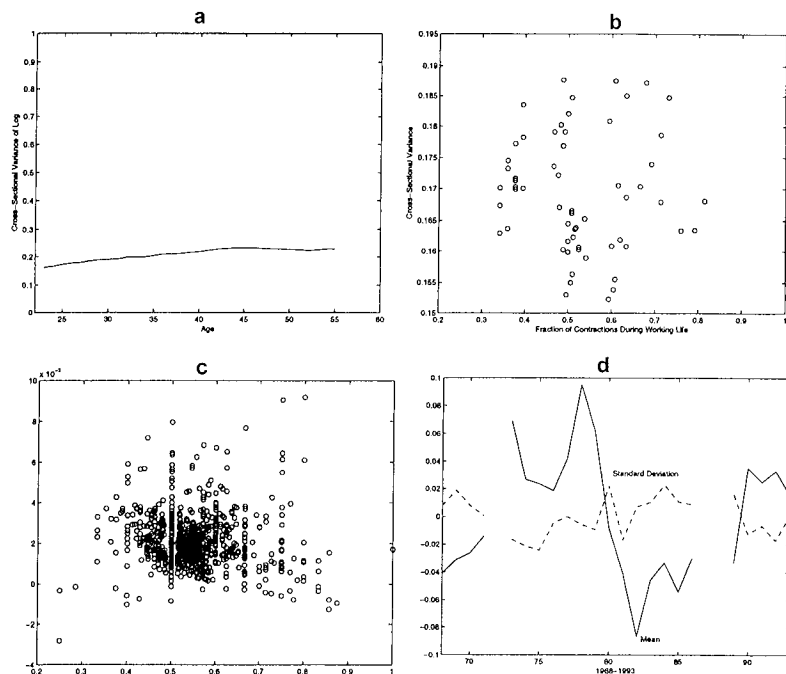


Fig. 2.—Cross-sectional variance of consumption, based on the idiosyncratic component (u_t^i from eq. [2]) of log consumption from the PSID, 1968–93, except the years 1972, 1987, and 1988, in which no consumption data are available. *a*, Cross-sectional variance by age. *b*, Macroeconomic history vs. cross-sectional variance (cohorts). *c*, Macroeconomic history vs. cross-sectional variance. *d*, Cross-sectional moments by time. See also the legend to fig. 1.

In summary, much of what we can infer graphically appears robust to the inclusion of information that is not represented in the graph. In particular, the age pattern in the cross-sectional variances (fig. 1*a*) suggests near-unit root behavior, and this is robust to the inclusion of autocovariances, the moments typically used to identify persistence.

A. Consumption Data

To this point we have focused on labor earnings. For many questions, most notably asset pricing, we are equally interested in consumption and how its cross-sectional distribution is related to aggregate variation. Figure 2 displays plots analogous to those in figure 1 using data on food expenditure from the PSID (the only consumption data available). Using these consumption data, table 3 replicates the estimation in table 2. Many of the patterns displayed in figure 1 for earnings seem to also appear for consumption. Figure 2*a* displays an almost linear rise in

TABLE 3
IDIOSYNCRATIC CONSUMPTION PROCESS: PARAMETER ESTIMATES

	ρ	σ_c	σ_E	σ_i	σ_a	p -Value
A. Exactly identified (no autocovariance)	.934 (.042)	.126 (.047)	.089 (.038)	.382 (.061)
B. Exactly identified (with autocovariances)	.911 (.037)	.160 (.058)	.115 (.042)	.278 (.058)	.269 (.053)	...
C. Overidentified, $H \times T$.862 (.015)	.222 (.059)	.172 (.043)	.283 (.040)	.267 (.013)	.999
D. Overidentified, NBER indicators	.893 (.021)	.172 (.073)	.133 (.071)	.288 (.043)	.246 (.013)	.999
E. Overidentified, unemployment indicators	.909 (.051)	.155 (.028)	.121 (.029)	.263 (.034)	.289 (.050)	.947
F. Overidentified, CRSP value-weighted return indicators	.913 (.050)	.127 (.016)	.095 (.023)	.295 (.005)	.303 (.057)	.999

NOTE.—Entries are GMM estimates, based on food consumption data from the PSID. We define a first-stage regression as in (1) except that the dependent variable is now log food consumption. We then employ a methodology identical to that used in table 2, where we treat the residual as the idiosyncratic consumption process. The PSID does not report consumption data for the years 1972, 1987, and 1988.

consumption dispersion. This leads to large autocorrelation estimates, although they are lower, ranging from 0.86 to 0.93. Consistent with the plots in figure 2, estimates of the fixed-effect variance are also lower, something we might expect given that we are limited to food data (i.e., overall inequality in earnings is likely to exceed that of food consumption). Finally, estimates of the persistent shock conditional variance still indicate countercyclical volatility. Further, consistent with figures 2*b* and *c*, these estimates are lower on average, and the difference between expansion and contraction is muted, with the peak to trough increase in the neighborhood of 30 percent instead of 75 percent. Further, the estimates are more sensitive to the definition of expansion/contraction. The extent to which these estimates are indicative of the behavior of a broader measure of consumption, obviously, remains to be seen.

VI. Temporal Variation in the Cross-Sectional Variance

Figure 1*d* indicates that, over the period 1968–93, the largest change in the age-pooled cross-sectional standard deviation of log earnings was about 10 percent. Similarly, the largest change from one year to the next was 4.4 percent. At first blush, this may seem inconsistent with our estimates, which indicate a change of roughly 75 percent over the course of the business cycle. The missing link is that the standard deviation in figure 1*d* is closely associated with the unconditional cross-sectional distribution of earnings, whereas the estimates in table 2 are associated with the conditional distribution. In this section we ask whether the two

can be quantitatively reconciled, thereby providing another corroborative check on our estimates.

The pooled cross-sectional variance in figure 1*d* is a weighted average of the conditional variances σ_h^2 and σ_c^2 . This is true in two senses. First, equation (3) indicates that, in the usual AR(1) sense, the cross-sectional variance among agents of a given age h is a moving average of the variances of past innovations. The only wrinkle is that the terms in the moving average change over time, depending on the macroeconomic history that the cohort has lived through. Second, the pooled variance is a weighted average of the cohort-specific variances, where the weights are the population shares. Algebraically,

$$v_t \equiv \tilde{\text{Var}}(u_{it}) = \sum_{h=1}^H \varphi_h \left\{ \sigma_\alpha^2 + \sigma_\epsilon^2 + \sum_{j=0}^{h-1} \rho^{2j} [I_{t-j} \sigma_E^2 + (1 - I_{t-j}) \sigma_C^2] \right\}, \quad (10)$$

where φ_h are population shares, and v_t denotes the pooled cross-sectional variance.

Inspection of equation (10) indicates that v_t is a moving average that varies between an upper and a lower bound. Given that $\sigma_c > \sigma_E$, the upper bound coincides with a long sequence of aggregate contractions (i.e., $I_{t-h} = 0$ for all $0 < h \leq H$) and is equal to $\sigma_\alpha^2 + \sigma_\epsilon^2 + [\sigma_C^2 / (1 - \rho^2)]$ for large H and $\rho < 1$. Similarly, the lower bound coincides with a long sequence of expansions and is equal to $\sigma_\alpha^2 + \sigma_\epsilon^2 + [\sigma_E^2 / (1 - \rho^2)]$. The way in which v_t varies between these extremes depends on persistence in both aggregate and idiosyncratic shocks and the economy's demographic structure (the φ_h 's). Figure 3 graphs this for one particular realization of the process (2), using U.S. data for the φ_h 's. It plainly illustrates the main point of this section, that large changes in the variance of the conditional distribution are necessarily associated with smaller changes in the pooled cross section.

We use this fact as another check on the plausibility of our estimates. We conduct a Monte Carlo experiment in which we obtain a large number of replications of the process (2). Each replication is a panel with a time dimension matching our PSID data (26 years) and a large number of individuals i in the cross section. We define two test statistics. The first is based on the difference between the maximum and the minimum values of v_t , which are observed in any given replication:

$$\max(\sqrt{v_t}) - \min(\sqrt{v_t}). \quad (11)$$

The second is based on the maximal absolute change between any two adjacent years:

$$\max(|\sqrt{v_t} - \sqrt{v_{t-1}}|). \quad (12)$$

We use the parameter estimates in row C of table 2, with aggregate




FIG. 3.—Simulated cross-sectional standard deviation. The solid line is one particular realization of the (population) pooled cross-sectional standard deviation of income shown in eq. (10), based on the process (2) from Sec. II. Parameter estimates are taken from row C of table 2. The upper line is the limit point, should all idiosyncratic shocks be drawn from the high-variance (contractionary) distribution. The lower line is the analogous lower limit point. Further details are provided in this section.

shocks following the transition matrix $\begin{bmatrix} \frac{2}{3} & \frac{1}{3} \\ \frac{1}{3} & \frac{2}{3} \end{bmatrix}$. We obtain the following results for the two test statistics. For the first, based on equation (11), the fraction of cases in which the maximal change across the 26 years exceeds the 10 percent value from figure 1*d* is 28 percent. For the second, based on equation (12), the fraction of cases in which the change from one year to the next exceeds 3 percent (3.5 percent) is 5 percent (< 1 percent).

In summary, the estimated large changes in the conditional variance reported in table 2 are quantitatively consistent with the relatively small changes we see in the pooled cross section, figure 1*d*. This serves to put the validity of our estimates on firmer ground, especially to the extent that variation in the pooled cross section is more easily and accurately measured.

VII. Last Thoughts

Our main finding is that PSID data on idiosyncratic labor earnings risk exhibit robust evidence of high persistence and countercyclical volatility.

We now conclude by comparing our estimates to those of related work and offering some economic interpretation.

We estimate an autocorrelation of roughly 0.95. Many other papers also advocate high persistence. MaCurdy (1982), Abowd and Card (1989), Carroll (1997), and Gourinchas and Parker (2002)—prominent papers on labor income dynamics—*restrict* labor income to be a unit root process. Hubbard et al. (1994) do not impose a unit root, but instead estimate an autocorrelation of 0.95. An exception is the paper by Heaton and Lucas (1996), who obtain estimates closer to 0.5.⁹ What distinguishes our work is an emphasis on age and its relation to cross-sectional variance. The linearly increasing pattern suggests near-unit root behavior. This is robust to the incorporation of what most other studies have emphasized: autocovariances. Our findings, therefore, represent new evidence in support of the Hubbard et al. estimates as well as the restrictions imposed by the other papers.

At face value, our estimates of transitory shocks seem large. We get an estimated standard deviation of roughly 0.25, which, in 2002 dollars, translates to an annual standard deviation of \$11,500 for the average worker in our panel, who earns just over \$45,000. While we cannot say how much of this represents measurement error and how much represents actual shocks, we can say that it is comparable to related work. Gourinchas and Parker (2002) and Hubbard et al. (1994), for example, obtain 0.21 and 0.17, respectively. Our estimates are somewhat higher, the most likely reason being additional measurement error in our relatively broad selection criterion.

We estimate the standard deviation of fixed effects to be roughly 0.40. The average young worker earns just over \$21,000 (in 2002 dollars), so this translates into inequality among the young with a standard deviation of roughly \$8,800. That this is plausible is apparent in figure 1*a*, where the standard deviation among the young is 0.55 (or \$12,480). We get a lower number here because (a) the transitory and persistent shocks also contribute to the initial variance (see eq. [3] with $h = 1$), and (b) the overidentifying moments in rows C and D of table 2 generate a reduction in overall variability, as the GMM estimator gives weight to moments not represented in figure 1*a*.

Our estimates of the conditional standard deviations of the persistent shock are roughly 0.21 in a contraction and 0.12 in an expansion. The frequency-weighted average is 0.17. Because the autocorrelation is almost unity, these are approximate estimates of the standard deviations

⁹ The main differences between our approach and Heaton and Lucas's is that they do not condition on age as we do, and they estimate one intercept term per household, whereas we make the stronger assumption that fixed effects are drawn from a single-parameter distribution. They also use the more traditional longitudinal panel, whereas we use the sequence of overlapping panels described in Sec. III.

of labor income *growth rates*. Equivalently, they translate into contraction/expansion conditional standard deviations of \$9,500 and \$5,300, based on the average worker's 2002 dollar earnings of \$45,000. By any measure, then, idiosyncratic risk is large. Related work reaches similar conclusions (relative to our average). Gourinchas and Parker (2002) estimate 0.17 for high school graduates and 0.15 for their overall sample. Carroll and Samwick (1997) obtain almost the same numbers. Hubbard et al. (1994) obtain 0.16 for high school graduates and 0.14 for college graduates. Our estimates, therefore, agree with related work, on average, and also suggest that cyclical variation is an important part of the story.

Finally, how large are the shocks relative to each other? We present a variance decomposition based on total lifetime uncertainty. Consider an unborn individual who is to receive lifetime labor income according to equation (1). Using a constant discount factor (for simplicity), we compute the variance of the present value of their future income (in levels, not logs):

$$\text{Var} \left[\sum_{h=23}^{62} 1.04^{-h} \exp(y_{it}^h) \right]. \quad (13)$$

We compute this variance by simulation.¹⁰ We do so three times, first with only the fixed-effect shocks set to zero, second with only the transitory shocks set to zero, and third with only the persistent shocks set to zero. We then compute the ratio of each of these three variances to the total variance in equation (13). We find that the fixed effects contribute 54.8 percent to the total, the persistent shocks contribute 44.5 percent, and the transitory shocks contribute 0.7 percent. The "size" of the fixed effects, then, is only slightly larger than that of the (relatively low conditional variance) persistent shocks, once the total working life is accounted for.

Appendix

Additional details regarding the data selection procedure, specific characteristics of the data, and the estimation procedure based on equations (1) and (7) are as follows.

Selection Criteria

A household is selected into a panel if the following conditions are met: (i) the head of the household is no younger than 22 and not older than 60, (ii) total

¹⁰ To simulate, we use the parametric specifications in eqq. (1), (2), and (6), the parameter estimates in tables 1 and 2 (row C), and a two-state Markov chain with transition matrix $\begin{bmatrix} \frac{2}{3} & \frac{1}{3} \\ \frac{1}{3} & \frac{2}{3} \end{bmatrix}$ for the process expansion/contraction needed to compute the indicator variables in eqq. (2).

earnings are positive in each of the sample years, (iii) total earnings growth rates are no larger than 20 and no lower than 1/20 in any consecutive years, and (iv) the household was not part of the Survey of Economic Opportunity.

Individual Data

The individual earnings data are based on the 1969–94 Family Files of the PSID. We use the PSID's Individual Files to track individuals across the different years of the Family Files. The definition of earnings includes wage earnings by the head of the household plus female wage earnings plus total transfers to the household. Total transfers include unemployment insurance, worker's compensation, transfers by nonhousehold family members, and several additional (minor) categories. Total earnings are then deflated to common 1968 dollars using the consumer price index. Earnings are then converted to rates per household member by dividing total earnings by family size.

Consumption data are based on total food consumption (the sum of expenditures of food for meals at home, meals outside the home, and the value of food stamps used). These questions were not available for panel years 1973, 1988, and 1989.

Three-Year Repeated Panels

Agents are selected into a panel if they meet the selection criteria above for the two years following the base year of the panel. We start with the 1969 PSID file (and therefore follow agents through PSID files from 1970 and 1971), which we denote panel 1968, since each PSID file provides information on the previous year's income. Our last three-year panel, which we denote as the 1991 panel, starts with the 1990 PSID file and ends with the 1994 PSID file. Hence, we end up with 24 three-year repeated panels. Table A1 reports summary statistics of the three-year repeated panels we use in our estimations.

Aggregate Data

The aggregate GNP levels used to define the indicator functions described in (7) were derived by merging GNP data from Gordon (1986) for the years 1910–58 with Citibase data for the years 1959–93. These levels were converted to 1968 dollars using the CPI and to per capita terms by dividing by total population figures from Citibase.

The NBER-based business cycles are derived from the monthly NBER definitions of contractions and expansions. We converted these monthly definitions into yearly ones. Our criterion, corresponding to the NBER definition, is to define years of recession as those years for which the majority of the months are contractionary. For cases in which contractions spanned less than 12 months but more than six months over two calendar years, we designated the first year as a recession. Our results are robust to close alternatives.

The "unemployment" definition of cycles is based on annual unemployment rates. This is constructed from two sources. For the years 1910–70 we used Historical Statistics of the United States (p. 135). For 1971–93, we used the Economic Report of the President (February 1999, p. 376). To convert these rates into contractionary years, we used information about the levels of unemployment as well as direction of changes. Recessions are defined as years for

TABLE A1
SUMMARY STATISTICS OF PANELS

YEAR	NUMBER OF HOUSEHOLDS	AGE OF HEAD HOUSEHOLD	
		Mean	Standard Deviation
1968	1,821	43.22	12.18
1969	1,923	42.83	12.39
1970	1,971	42.27	12.63
1971	2,036	41.57	12.69
1972	2,088	41.22	12.61
1973	2,124	40.75	12.64
1974	2,216	40.87	12.73
1975	2,282	40.57	12.79
1976	2,337	40.22	12.76
1977	2,406	40.12	12.70
1978	2,452	40.18	12.59
1979	2,535	39.75	12.52
1980	2,565	39.76	12.49
1981	2,584	39.82	12.36
1982	2,619	39.75	12.19
1983	2,646	39.99	12.02
1984	2,728	39.93	11.99
1985	2,768	40.12	12.07
1986	2,778	40.38	11.83
1987	2,819	40.55	11.65
1988	2,819	40.90	11.44
1989	2,856	40.73	11.29
1990	2,806	40.78	11.18
1991	2,610	41.80	11.23

NOTE.—The entries correspond to three-year panels. The entries under the year column correspond to the base year of each panel. The age statistics correspond to the head of the household.

which the unemployment rate was greater than 7.5 percent. If the rate was greater than 7.5 percent but fell more than 3 percent relative to the previous year, it was not counted as a recession year. In addition, years for which the unemployment rate had risen by more than 3 percent were defined as recession years.

The "returns" definition is based on the real (CPI-deflated) value-weighted return from the Center for Research in Security Prices (CRSP). We define contractionary years to be ones in which the real return is negative.

Estimation Procedure

Our estimation procedure has two distinct steps. In the first stage we estimate equation (1). In the second stage we estimate the system given in (7).

To recover u_{it}^h in the first stage, we need to specify $g(Y_t, x_{it}^h)$. We let $g(\cdot)$ be represented by a simple linear regression using as regressors year dummies (corresponding to Y_t) and age, age squared, age cubed, and education (corresponding to individual-specific attributes x_{it}^h). Let the parameters of this first-stage regression be summarized by θ_1 . That is, $\psi_1(y_{i,t}, Y_t, x_{it}^h, \theta_1) \equiv y_{it}^h - \theta_1[1, Y_t,$

x_{it}^h]. All the coefficients are significant, and the R^2 is .23. These are displayed in table 1.

Next, let the parameters of the system in (7) be denoted by θ_2 (i.e., ρ , σ_ϵ , σ_α , σ_C , σ_E , and σ_α). The joint system we estimate can be written compactly as

$$E \begin{bmatrix} \psi_1(y_{it}^h, Y_{it}^h, x_{it}^h, \theta_1) \\ \psi_2(y_{it}^h, Y_{it}^h, \theta_1, \theta_2) \end{bmatrix} = 0, \quad (A1)$$

where ψ_1 and ψ_2 are the moment conditions corresponding to (1) and (7), respectively. The triangular structure of the moment condition allows us to get consistent estimates of θ_1 using only ψ_1 . The weighting matrix is trivially the identity since this is always an exactly identified system. We then estimate θ_2 using moment conditions ψ_2 . This second step incorporates the standard errors in estimating θ_1 using the standard two-step GMM procedure as in Ogaki (1993).

There are two additional complications that arise in our setup: the overlapping structure of our repeated panels when we use autocovariances and the non-balanced panel.

For each moment condition based on, say, panel year t , an MA(2) correction is added to the estimate of the covariance matrix associated with moment conditions ψ_2 . Specifically, define the empirical residuals of the moment conditions ψ_2 to be

$$\begin{aligned} \psi_{2,i,t}^{h,0} &\equiv \left((u_{i,t}^{h,0})^2 - \left\{ \sigma_\alpha^2 + \sigma_\epsilon^2 + \sum_{j=0}^{h-1} \rho^{2j} [I_{t-j} \sigma_C^2 + (1 - I_{t-j}) \sigma_E^2] \right\} \right), \\ \psi_{2,i,t}^{h,1} &\equiv \left(u_{i,t}^{h,1} u_{i,t+1}^{h+1,1} - \left\{ \sigma_\alpha^2 + \rho \sum_{j=1}^{h-1} \rho^{2(j-1)} [I_{t-j} \sigma_C^2 + (1 - I_{t-j}) \sigma_E^2] \right\} \right), \\ \psi_{2,i,t}^{h,2} &\equiv \left(u_{i,t}^{h,2} u_{i,t+2}^{h+2,2} - \left\{ \sigma_\alpha^2 + \rho^2 \sum_{j=1}^{h-1} \rho^{2(j-2)} [I_{t-j} \sigma_C^2 + (1 - I_{t-j}) \sigma_E^2] \right\} \right), \end{aligned}$$

where the superscript t in u denotes the base year of the panel from which this agent is selected. The superscript j denotes whether the moment is based on variances ($j = 0$) or first and second autocovariances ($j = 1, 2$). By assumption, $\psi_{2,i,t}^{h,j}$ is not correlated with $\psi_{2,i,t+k}^{h,l}$ for all $k \neq 0$ and $j, l = 0, 1, 2$. It can be easily shown that, because of the overlap of the sample, for each t , $\psi_{2,i,t}^{h,0}$, $\psi_{2,i,t}^{h,1}$, and $\psi_{2,i,t}^{h,2}$ are correlated. We stack the repeated three moment conditions and use sample counterparts to estimate these covariance terms; the covariance matrix is block-diagonal, and each 3×3 block has nonempty off-diagonal elements. In practice, our results do not seem to be sensitive to this correction.

Because of the block-diagonal structure described above, the nonbalanced panel allows for the construction of standard GMM statistics. Let N^* be the minimum sample size across all moments, and let other sample moments, say moment j , be of size κ_j times N^* . The standard GMM asymptotics follow from assuming that sample sizes grow in these ratios. The covariance matrix used is then constructed using each sample moment scaled by its κ_j .

References

- Abowd, John M., and David Card. 1989. "On the Covariance Structure of Earnings and Hours Changes." *Econometrica* 57 (March): 411-45.

- Alvarez, Fernando, and Urban J. Jermann. 2001. "Quantitative Asset Pricing Implications of Endogenous Solvency Constraints." *Rev. Financial Studies* 14 (Winter): 1117-51.
- Balduzzi, Pierluigi, and Tong Yao. 2000. "Does Heterogeneity Matter for Asset Pricing?" Manuscript. Boston: Boston Coll.
- Bernanke, Ben, and Mark Gertler. 1989. "Agency Costs, Net Worth, and Business Fluctuations." *A.E.R.* 79 (March): 14-31.
- Boeri, Tito. 1996. "Is Job Turnover Countercyclical?" *J. Labor Econ.* 14 (October): 603-25.
- Bray, Alon, George M. Constantinides, and Christopher C. Geczy. 2002. "Asset Pricing with Heterogeneous Consumers and Limited Participation: Empirical Evidence." *J.P.E.* 110 (August): 793-824.
- Caballero, Ricardo J., and Mohamad L. Hammour. 1994. "The Cleansing Effect of Recessions." *A.E.R.* 84 (December): 1350-68.
- Carlstrom, Charles T., and Timothy S. Fuerst. 1997. "Agency Costs, Net Worth, and Business Fluctuations: A Computable General Equilibrium Analysis." *A.E.R.* 87 (December): 893-910.
- Carroll, Christopher D. 1997. "Buffer-Stock Saving and the Life Cycle/Permanent Income Hypothesis." *Q.J.E.* 112 (February): 1-55.
- Carroll, Christopher D., and Andrew A. Samwick. 1997. "The Nature of Precautionary Wealth." *J. Monetary Econ.* 40 (September): 41-71.
- Cogley, Timothy. 2002. "Idiosyncratic Risk and the Equity Premium: Evidence from the Consumer Expenditure Survey." *J. Monetary Econ.* 49 (March): 309-34.
- Constantinides, George M., and Darrell Duffie. 1996. "Asset Pricing with Heterogeneous Consumers." *J.P.E.* 104 (April): 219-40.
- Cooley, Thomas, Ramon Marimon, and Vincenzo Quadrini. 2002. "Aggregate Consequences of Limited Contract Enforceability." Manuscript. New York: New York Univ.
- Davis, Steven J., and John C. Haltiwanger. 1992. "Gross Job Creation, Gross Job Destruction, and Employment Reallocation." *Q.J.E.* 107 (August): 819-63.
- Deaton, Angus, and Christina Paxson. 1994. "Intertemporal Choice and Inequality." *J.P.E.* 102 (June): 437-67.
- Foote, Christopher L. 1998. "Trend Employment Growth and the Bunching of Job Creation and Destruction." *Q.J.E.* 113 (August): 809-34.
- Gordon, Robert J., ed. 1986. *The American Business Cycle: Continuity and Change*. Chicago: Univ. Chicago Press.
- Gourinchas, Pierre-Olivier, and Jonathan A. Parker. 2002. "Consumption over the Life Cycle." *Econometrica* 70 (January): 47-89.
- Heaton, John, and Deborah J. Lucas. 1996. "Evaluating the Effects of Incomplete Markets on Risk Sharing and Asset Pricing." *J.P.E.* 104 (June): 443-87.
- Hubbard, R. Glenn, Jonathan Skinner, and Stephen P. Zeldes. 1994. "The Importance of Precautionary Motives in Explaining Individual and Aggregate Saving." *Carnegie-Rochester Conf. Ser. Public Policy* 40 (June): 59-125.
- Krebs, Tom. 2002. "Growth and Welfare Effects of Business Cycles in Economies with Idiosyncratic Human Capital Risk." Manuscript. Providence, R.I.: Brown Univ.
- . 2003. "Human Capital Risk and Economic Growth." *Q.J.E.* 118 (May): 709-44.
- Krusell, Per, and Anthony A. Smith, Jr. 1997. "Income and Wealth Heterogeneity, Portfolio Choice, and Equilibrium Asset Returns." *Macroeconomic Dynamics* 1 (2): 387-422.

- . 1999. "On the Welfare Effects of Eliminating Business Cycles." *Rev. Econ. Dynamics* 2 (January): 245–72.
- Lucas, Deborah J. 1994. "Asset Pricing with Undiversifiable Risk and Short Sales Constraints: Deepening the Equity Premium Puzzle." *J. Monetary Econ.* 34 (December): 325–41.
- Lustig, Hanno. 2001. "The Market Price of Aggregate Risk and the Wealth Distribution." Manuscript. Stanford, Calif.: Stanford Univ.
- MacCurdy, Thomas E. 1982. "The Use of Time Series Processes to Model the Error Structure of Earnings in a Longitudinal Data Analysis." *J. Econometrics* 18 (January): 83–114.
- Mankiw, N. Gregory. 1986. "The Equity Premium and the Concentration of Aggregate Shocks." *J. Financial Econ.* 17 (September): 211–19.
- Marcet, Albert, and Kenneth J. Singleton. 1999. "Equilibrium Asset Prices and Savings of Heterogeneous Agents in the Presence of Incomplete Markets and Portfolio Constraints." *Macroeconomic Dynamics* 3 (June): 243–77.
- Ogaki, Masao. 1993. "Generalized Method of Moments: Econometric Applications." In *Econometrics, Handbook of Statistics*, vol. 11, edited by G. S Maddala, C. R. Rao, and H. D. Vinod. Amsterdam: Elsevier Sci.
- Rampini, Adriano. 2000. "Entrepreneurial Activity, Risk and the Business Cycle." Manuscript. Evanston, Ill.: Northwestern Univ.
- Sarkissian, Sergei. 2003. "Incomplete Consumption Risk Sharing and Currency Risk Premiums." *Rev. Financial Studies* 16 (Fall): 983–1005.
- Storesletten, Kjetil, Chris I. Telmer, and Amir Yaron. 2001a. "Asset Pricing with Idiosyncratic Risk and Overlapping Generations." Manuscript. Pittsburgh: Carnegie Mellon Univ.
- . 2001b. "The Welfare Cost of Business Cycles Revisited: Finite Lives and Cyclical Variation in Idiosyncratic Risk." *European Econ. Rev.* 45 (June): 1311–39.
- . In press. "Consumption and Risk Sharing over the Life Cycle." *J. Monetary Econ.*
- White, Halbert. 1980. "A Heteroskedasticity-Consistent Covariance Matrix Estimator and a Direct Test for Heteroskedasticity." *Econometrica* 48 (May): 817–38.